

Lecture 3: Data Cleaning & Wrangling with Tidyverse

James Sears*

AFRE 891/991 FS 26

Michigan State University

*Parts of these slides are adapted from [“Data Science for Economists”](#) by Grant McDermott.

Table of Contents

1. Prologue
2. Tidyverse Overview
3. Pipes
4. Data Wrangling with dplyr
 - filter, arrange, select, mutate, summarise, joins
5. Data Tidying with tidyr
 - pivot longer / pivot wider
 - separate
 - unite
6. Summary

Prologue

What is "tidy" data?

Resources:

- [Vignette](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)
- [Online Book: R 4 Data Science](#)

Key points:

1. Each **variable** forms a **column**.
2. Each **observation** forms a **row**.
3. Each **type of observational unit** forms a **table**.

Basically, tidy data is more likely to be [long \(i.e. narrow\) format](#) than wide format.

Checklist

R packages you'll need today

☒ tidyverse

☒ nycflights13

Let's hold off on loading/installing these right now

Tidyverse Overview

Tidyverse vs. base R

One thing to note before we dive into **tidyverse**: there is often a **direct correspondence** between a tidyverse command and its base R equivalent.

These generally follow a `tidyverse::snake_case` VS `base::period.case` rule.
E.g. Compare:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

If you call up the above examples, you'll see that the tidyverse alternative typically offers some enhancements or other useful options (and sometimes restrictions) over its base counterpart.

Tidyverse vs. base R

One thing to note before we dive into **tidyverse**: there is often a **direct correspondence** between a tidyverse command and its base R equivalent.

These generally follow a `tidyverse::snake_case` VS `base::period.case` rule.
E.g. Compare:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

Remember: There are (almost) always multiple ways to achieve a single goal in R.

Tidyverse Packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We see that we have actually loaded a number of packages (which could also be loaded individually): **ggplot2**, **tibble**, **dplyr**, etc.

- We can also see information about the package versions and some **namespace conflicts**.

Tidyverse Packages

The tidyverse actually comes with a **lot more packages** that aren't loaded automatically.

```
tidyverse_packages()
```

```
## [1] "broom"          "conflicted"     "cli"            "dbplyr"
## [5] "dplyr"          "dtplyr"         "forcats"        "ggplot2"
## [9] "googledrive"    "googlesheets4" "haven"          "hms"
## [13] "httr"           "jsonlite"       "lubridate"      "magrittr"
## [17] "modelr"         "pillar"         "purrr"          "ragg"
## [21] "readr"          "readxl"         "reprex"         "rlang"
## [25] "rstudioapi"     "rvest"          "stringr"        "tibble"
## [29] "tidyr"          "xml2"           "tidyverse"
```

We'll use several of these additional packages during the remainder of this course: **haven** for loading STATA files, **lubridate** for working with dates, **rvest** package for webscraping.

Tidyverse Packages

This week we're going to focus on two packages:

1. **dplyr**
2. **tidyr**

These are the workhorse packages for cleaning and wrangling data that you will likely make the most use of (alongside **ggplot2**, which we'll cover more in depth later on)

- Data cleaning and wrangling occupies an **inordinate amount of time**, no matter where you are in your research career.

Pipes



Pipes

The tidyverse **pipe operator** `%>%` is one of its **coolest features**.

- It lets you send (i.e. "pipe") intermediate output to another command.
 - Automatically passes what's "upstream" as the first argument of the "downstream" function
- In other words, it allows us to chain together a sequence of simple operations and thereby implement a more complex operation while preserving legibility
 - Avoids nesting multiple functions, creating many intermediate objects

Let's demonstrate with an example.

Pipes

Suppose you're a big German car fan and want to see average fuel efficiency of their models for 1999-2008. The discrete operations involved are

1. Load the dataset (`mpg` "loaded" by tidyverse)
2. Filter the data to Audi and Volkswagen (`filter()`)
3. Group the data by model (`group_by()`)
4. Summarise average highway mileage (`summarise()`)

Without pipes, we would need to

1. Assign/reassign intermediate objects to memory after each step (repetitive), or
2. Nest a lot of functions (hard to read)

Without Pipes: 1. Assign/Reassign

Alternatively, we could **assign/reassign** and create **intermediate objects**.

```
cars ← mpg
german_cars ← filter(mpg, manufacturer %in% c("audi", "volkswagen"))
unique(cars$manufacturer)
german_cars_grp ← group_by(german_cars, manufacturer)
summarise(german_cars_grp,
          hwy_mean = mean(hwy),
          hwy_max = max(hwy),
          "Minimum Highway MPG" = min(hwy)
        )

rm(german_cars, german_cars_grp)
```

This is stac-ca-to to read and leaves us with a bunch of intermediate objects that we'll need to deal with.

```
rm(cars, german_cars, german_cars_grp)
```

Without Pipes: 2. Nest

The **nested approach** is harder to read and **totally inverts the logical order**

```
summarise(group_by(filter(mpg, manufacturer %in% c("audi",  
"volkswagen")), manufacturer, model),  
          hwy_mean = mean(hwy))
```

- The final operation comes first!
- Who wants to read things inside out??

With Pipes

The below line does exactly the same thing through the power of pipes:

```
mpg %>%  
  filter(manufacturer %in% c("audi", "volkswagen")) %>%  
  group_by(manufacturer, model) %>%  
  summarise(hwy_mean = mean(hwy))
```

With pipes the line reads from left to right, exactly how I thought of the operations in my head.

- Take this object (`mpg`), do this (`filter`), then do this (`group_by`), etc.

Pipes: Improved Readability

The piped version of the code is **even more readable** if we write it **over several lines**. Here it is again and, this time, I'll run it so we can see the output:

```
pacman::p_load(haven)

mpg %>%
  filter(manufacturer %in% c("audi", "volkswagen")) %>%
  group_by(manufacturer, model) %>%
  summarise(hwy_mean = mean(hwy))
```



```
## # A tibble: 7 × 3
## # Groups:   manufacturer [2]
##   manufacturer model      hwy_mean
##   <chr>          <chr>      <dbl>
## 1 audi          a4          28.3
## 2 audi          a4 quattro  25.8
## 3 audi          a6 quattro  24
## 4 volkswagen    gti         27.4
## 5 volkswagen    jetta       29.1
```

Pipes: Improved Readability

```
mpg %>%  
  filter(manufacturer %in% c("audi", "volkswagen")) %>%  
  group_by(manufacturer, model) %>%  
  summarise(hwy_mean = mean(hwy))
```

At each line, the **upstream object/output** (i.e. the `mpg` df) is being passed into the **downstream function** (i.e. `filter()`) as the first argument.

Remember: Using vertical space **costs nothing** and makes for much more readable/writeable code than cramming things horizontally.

PS — The pipe is originally from the [magrittr](#) package (hence the [not-a-pipe image](#)) earlier, which can do some other cool things if you're inclined to explore.

Pipes: Alternate Argument

By default, the pipe sends what's upstream into the **first argument** of the downstream function.

If you want to pass it to a **different argument**, place `.` in the target argument.

Default: $\ln(5)$

```
5 %>%  
  log()
```

```
## [1] 1.609438
```

Different Argument: $\log_5(10)$

```
5 %>%  
  log(10, .)
```

```
## [1] 1.430677
```

The base R pipe: |>

The magrittr pipe has proven so successful and popular, that as of R 4.1.0 the R core team added a "native" pipe, denoted `▷`.¹ For example:

```
mtcars ▷ subset(cyl=4) ▷ head()
```

¹ That's actually a `|` followed by a `>`. The default font on these slides just makes it look extra fancy.

The base R pipe: different arguments

One **difference** in base R pipe syntax: piping into **different arguments**:

- Use `_` (underscore) in place of `.`
- Must **name the argument** you're piping into

```
5 ▷  
log(10, _)
```

```
## Error in log(10, "_"): pipe  
placeholder can only be used as a  
named argument (<input>:2:3)
```

```
5 ▷  
log(10, base = _)
```

```
## [1] 1.430677
```

dplyr

Aside: Updating Packages

- Please make sure that you are running **at least dplyr 1.0.0** before continuing.
- 1.0.0 has been around for a while now (currently on 1.2.1), but if you have an old old version of dplyr, these functions won't be available
- As well, it's a good idea to frequently update all your packages!

```
packageVersion('dplyr')
```

```
## [1] '1.2.1'
```

```
# install.packages('dplyr') ## install updated version if < 1.0.0
```

Note: You can turn off **dplyr's** notifications about grouping variables with

```
options(dplyr.summarise.inform = FALSE) ## Add to .Rprofile to make permanent
```

Key dplyr Verbs

There are **five key dplyr verbs** that you need to learn.

1. `filter()`: Filter (i.e. subset) rows based on their values.
2. `arrange()`: Arrange (i.e. reorder) rows based on their values.
3. `select()`: Select (i.e. subset or arrange) columns by their names:
4. `mutate()`: Create new columns or modify existing columns.
5. `summarise()`: Collapse multiple rows into a single summary value.²

Let's practice these commands together using the `starwars` data frame that comes pre-packaged with dplyr.

² `summarize` with a "z" works too, R doesn't discriminate.

1) dplyr::filter

`filter(df, condition1, ...)` only keeps rows that match the specified condition(s)

```
View(starwars)
```

Character variable: (in)equality,
regular expressions

```
starwars ▶  
filter( species = "Human")
```

Numeric variable: (in)equality,
comparisons

```
starwars %>%  
filter(  
  between(mass, 75, 150)  
)
```

1) dplyr::filter

`filter()` only keeps rows that match the specified condition(s)

```
starwars ►  
filter( species = "Human")
```

```
## # A tibble: 35 × 14
```

```
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender
```

```
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>
```

```
## 1 Luke Sk...    172    77 blond      fair       blue       19  
male masculin...
```

```
## 2 Darth V...    202   136 none       white      yellow     41.9  
male masculin...
```

```
## 3 Leia Or...    150    49 brown      light      brown      19  
female feminin...
```

```
## 4 Owen La...    178   120 brown, gr... light      blue       52  
male masculin...
```

```
## 5 Beru Wh...    165    75 brown      light      blue       47  
female feminin...
```

1) dplyr::filter

We can chain multiple filter commands with the pipe (`%>%`, `▷`), or supply them as separate arguments within a single filter command.

```
starwars ▷  
filter(species = "Human", mass > 75, mass < 85,  
       height %in% 170:180)
```

```
## # A tibble: 5 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1 Luke Sky...   172    77 blond      fair       blue        19  
male  mascu...  
## 2 Han Solo     180    80 brown      fair       brown        29  
male  mascu...  
## 3 Wedge An...  170    77 brown      fair       hazel        21  
male  mascu...  
## 4 Lando Ca...  177    79 black      dark       brown        31  
male  mascu...
```

1) dplyr::filter

Regular expressions work well too. Using Base R:

```
starwars %>%  
  filter(grepl("Skywalker", name))  
  
## # A tibble: 3 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1 Luke Sky...   172    77 blond      fair       blue        19  
male  mascu...  
## 2 Anakin S...   188    84 blond      fair       blue        41.9  
male  mascu...  
## 3 Shmi Sky...   163    NA black      fair       brown        72  
fema... femin...  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

1) dplyr::filter

Or **stringr** functions (the tidyverse string package)

- Note the different argument order

```
starwars <- starwars %>%  
  mutate(skywalker = str_detect(name, "Skywalker"),  
         not_skywalker = !str_detect(name, "Skywalker"))  
  
sky_check <- select(starwars, name, skywalker) %>%  
  arrange(desc(skywalker), name)
```

Identifying Missing Data

A very common `filter()` use case is identifying (or removing) observations with **missing values**

```
filter(starwars, is.na((height))) %>%  
  nrow()
```

```
## [1] 6
```

To **remove missing observations**, simply use negation:

```
filter(!is.na(height))
```

```
filter(starwars, !is.na((height))) %>%  
  nrow()
```

```
## [1] 81
```

Removing Missing Data

Or use the convenient `drop_na()` verb:

- Drops `NA` on the named variables
- If none specified, drops `NA` in **any variable**

```
nrow(starwars)
```

```
## [1] 87
```

```
starwars %>%  
  drop_na(height) %>%  
  nrow()
```

```
## [1] 81
```

Removing Missing Data

Or use the convenient `drop_na()` verb:

- Drops `NA` on the named variables
- If none specified, drops `NA` in **any variable**

```
nrow(starwars)
```

```
## [1] 87
```

```
starwars %>%  
  drop_na() %>%  
  nrow()
```

```
## [1] 29
```

2) dplyr::arrange

`arrange()` **arranges/sorts rows** based on values of a variable/variables

- **numeric:** ascending order
- **character:** alphabetically (try this on `name` variable)

```
starwars %>%  
  arrange(birth_year) %>%  
  head()
```

```
## # A tibble: 6 × 16  
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1 Wicket S...    88    20 brown      brown      brown          8  
male  mascu...  
## 2 IG-88         200   140 none       metal      red           15  
none  mascu...  
## 3 Luke Sky...   172    77 blond      fair       blue           19  
male  mascu...
```

2) dplyr::arrange

We can also arrange items in **descending order** using `desc()`.

```
starwars %>%  
  arrange(desc(birth_year))
```

```
## # A tibble: 87 × 16  
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1 Yoda          66    17 white      green      brown      896  
male  masculi...  
## 2 Jabba D...    175   1358 <NA>      green-tan... orange      600  
herm... masculi...  
## 3 Chewbac...    228    112 brown      unknown    blue       200  
male  masculi...  
## 4 C-3PO        167     75 <NA>      gold       yellow     112  
none  masculi...  
## 5 Dooku         193     80 white      fair       brown     102  
male  masculi...
```

2) dplyr::arrange

We can also **nested sort** by including **multiple variables**

- Sort on first variable, then ties on the next, etc.

```
sw_ar_1 <- starwars %>%  
  arrange(desc(birth_year), name)
```

3) dplyr::select

`select()` lets you **subset columns by name/condition**

- Places variable name arguments in specified order
- Can also use `first:last` syntax for currently consecutive columns

```
starwars %>%  
  select(  
    name, species, height:mass  
  ) %>% head()
```

```
## # A tibble: 6 × 4  
##   name          species height  mass  
##   <chr>         <chr>    <int> <dbl>  
## 1 Luke Skywalker Human      172    77  
## 2 C-3PO         Droid     167    75  
## 3 R2-D2         Droid      96    32  
## 4 Darth Vader   Human     202   136  
## 5 Leia Organa   Human     150    49  
## 6 Owen Lars     Human     178   120
```

3) dplyr::select

`select()` lets you **subset columns by name/condition**

- Places variable name arguments in specified order
- Can also use `first:last` syntax for currently consecutive columns
- Deselect columns with `-` in front of the name(s).

```
starwars %>%  
  select(  
    name:sex, -mass  
  ) %>% head()
```

```
## # A tibble: 6 × 7
```

```
##   name          height hair_color  skin_color  eye_color birth_year
```

```
sex
```

```
##   <chr>          <int> <chr>      <chr>      <chr>      <dbl>
```

```
<chr>
```

```
## 1 Luke Skywalker    172 blond      fair        blue         19
```

```
male
```

```
## 2 C-3PO             167 <NA>      gold        yellow        112  38 / 110
```

```
none
```

3) dplyr::select

`select()` also works well with [helper functions](#)

```
starwars %>%  
  select(  
    where(is.logical)  
  )
```

```
## # A tibble: 87 × 2  
##   skywalker not_skywalker  
##   <lgl>      <lgl>  
## 1 TRUE      FALSE  
## 2 FALSE     TRUE  
## 3 FALSE     TRUE  
## 4 FALSE     TRUE  
## 5 FALSE     TRUE  
## 6 FALSE     TRUE  
## 7 FALSE     TRUE  
## 8 FALSE     TRUE  
## 9 FALSE     TRUE  
## 10 FALSE    TRUE
```

3) dplyr::select

`select()` also works well with [selection helper functions](#)

```
starwars %>%  
  select(  
    ends_with("color")  
  )
```

```
## # A tibble: 87 × 3  
##   hair_color    skin_color eye_color  
##   <chr>         <chr>      <chr>  
## 1 blond        fair       blue  
## 2 <NA>         gold       yellow  
## 3 <NA>         white, blue red  
## 4 none         white      yellow  
## 5 brown        light      brown  
## 6 brown, grey  light      blue  
## 7 brown        light      blue  
## 8 <NA>         white, red red  
## 9 black        light      brown  
## 10 auburn, white fair       blue-gray
```

3) dplyr::select

`select()` also works well with selection helper functions

```
starwars %>%  
  select(name, contains("color"))
```

```
## # A tibble: 87 × 4  
##   name                hair_color    skin_color eye_color  
##   <chr>              <chr>        <chr>      <chr>  
## 1 Luke Skywalker    blond        fair        blue  
## 2 C-3PO             <NA>        gold        yellow  
## 3 R2-D2             <NA>        white, blue red  
## 4 Darth Vader       none        white        yellow  
## 5 Leia Organa       brown        light        brown  
## 6 Owen Lars         brown, grey  light        blue  
## 7 Beru Whitesun Lars brown        light        blue  
## 8 R5-D4             <NA>        white, red   red  
## 9 Biggs Darklighter black        light        brown  
## 10 Obi-Wan Kenobi    auburn, white fair        blue-gray  
## # i 77 more rows
```

3) dplyr::select

You can also **rename** some (or all) of your selected variables while changing places.

```
starwars %>%  
  select(height_cm = height, everything()) %>%  
  head()
```

```
## # A tibble: 6 × 16  
##   height_cm name      mass hair_color skin_color eye_color birth_year sex  
gender  
##       <int> <chr>   <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1       172 Luke ...    77 blond      fair        blue        19  
male masculin  
## 2       167 C-3PO      75 <NA>      gold        yellow      112  
none masculin  
## 3        96 R2-D2      32 <NA>      white, bl... red         33  
none masculin  
## 4       202 Darth...  136 none      white       yellow      41.9  
male masculin
```

dplyr::rename

Note: you can **rename columns** without subsetting with `rename`.

3) everything()

The `select( ..., everything())` option is another useful shortcut if you only want to bring some variable(s) to the "front" of a data frame.

```
starwars %>%  
  select(species, homeworld, everything()) %>%  
  head(5)
```

```
## # A tibble: 5 × 16  
##   species homeworld name          height  mass hair_color skin_color  
eye_color  
##   <chr>    <chr>    <chr>          <int> <dbl> <chr>      <chr>  
<chr>  
## 1 Human   Tatooine  Luke Skywalker   172    77 blond      fair  
blue  
## 2 Droid   Tatooine  C-3PO            167    75 <NA>      gold  
yellow  
## 3 Droid   Naboo     R2-D2             96    32 <NA>      white, blue  
red  
## 4 Human   Tatooine  Darth Vader      202   136 none      white  
yellow
```

3) dplyr::select

Some other selection helper functions:

Operator/Function	Role
<code>!</code>	Take complement of variable set
<code>&</code>	intersection of two sets
<code>last_col()</code>	furthest right column
<code>starts_with("D")</code>	names starting with "D"
<code>matches("^[a-d]")</code>	names matching regular expression "^[a-d]"
<code>num_range(x, 1:10)</code>	names following pattern <code>x1</code> , <code>x2</code> , ..., <code>x10</code>
<code>all_of(vars)</code> / <code>any_of(vars)</code>	matches names stored in character vector <code>vars</code>
<code>group_cols()</code>	variables defining group structure

Aside: dplyr::relocate

Note that the function `relocate()` uses the same syntax as `select()` to move **groups of columns at once**.

Add variables separated by commas to move them **to the front**

```
starwars %>%
  relocate(
    ends_with("_color"), homeworld
  ) %>%
  head()
```

A tibble: 6 × 16

	hair_color	skin_color	eye_color	homeworld	name	height	mass	birth_year	sex
##	<chr>	<chr>	<chr>	<chr>	<chr>	<int>	<dbl>	<dbl>	<chr>
## 1	blond	fair	blue	Tatooine	Luke...	172	77	19	male
## 2	<NA>	gold	yellow	Tatooine	C-3PO	167	75	112	none

Aside: dplyr::relocate

Can also use arguments `.after` / `.before` to place the column(s) in **specific locations**

```
starwars %>%  
  relocate(  
    species,  
    .before = height  
  ) %>%  
  head()
```

```
## # A tibble: 6 × 16  
##   name      species height  mass hair_color skin_color eye_color  
birth_year sex  
##   <chr>      <chr>    <int> <dbl> <chr>      <chr>      <chr>  
<dbl> <chr>  
## 1 Luke Sk... Human      172    77 blond      fair       blue  
19   male  
## 2 C-3PO      Droid      167    75 <NA>      gold       yellow  
112  none  
## 3 R2-D2      Droid       96    32 <NA>      white, bl... red
```

4) dplyr::mutate

You can use `mutate()` to **create new columns** from scratch, or (more commonly) **transform existing columns**.

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(dog_years = birth_year * 7) %>%  
  head()
```

```
## # A tibble: 6 × 3  
##   name          birth_year dog_years  
##   <chr>         <dbl>     <dbl>  
## 1 Luke Skywalker      19        133  
## 2 C-3PO              112        784  
## 3 R2-D2               33        231  
## 4 Darth Vader        41.9       293.  
## 5 Leia Organa        19        133  
## 6 Owen Lars          52        364
```

4) dplyr::mutate

Note: `mutate()` is **order aware**, so you can chain multiple mutates in a single call.

```
starwars %>%
  select(name, birth_year) %>%
  mutate(
    dog_years = birth_year * 7, ## Separate with a comma
    comment = paste0(name, " is ", dog_years, " in dog years.")
  ) %>% head()
```

```
## # A tibble: 6 × 4
##   name          birth_year dog_years comment
##   <chr>          <dbl>     <dbl> <chr>
## 1 Luke Skywalker      19        133 Luke Skywalker is 133 in dog
years.
## 2 C-3PO              112        784 C-3PO is 784 in dog years.
## 3 R2-D2              33        231 R2-D2 is 231 in dog years.
## 4 Darth Vader       41.9       293.3 Darth Vader is 293.3 in dog
years.
## 5 Leia Organa       19         133 Leia Organa is 133 in dog years.
```

4) dplyr::mutate

Boolean, logical and conditional operators all work well with `mutate()` too.

```
starwars %>%  
  select(name, height) %>%  
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) %>%  
  mutate(tall1 = height > 180) %>%  
  mutate(tall2 = if_else(height > 180, "Tall", "Short")) ## Same effect,  
  but can choose labels
```

```
## # A tibble: 2 × 4  
##   name          height tall1 tall2  
##   <chr>         <int> <lgl> <chr>  
## 1 Luke Skywalker    172 FALSE Short  
## 2 Anakin Skywalker   188  TRUE  Tall
```

4) dplyr::mutate

Lastly, combining `mutate()` with the recent `across()` feature³ allows you to easily work on a **subset of variables**:

```
starwars %>%  
  select(name:eye_color) %>%  
  mutate(across(where(is.character), toupper)) %>%  
  head(5)
```

```
## # A tibble: 5 × 6  
##   name          height  mass hair_color skin_color eye_color  
##   <chr>         <int> <dbl> <chr>      <chr>      <chr>  
## 1 LUKE SKYWALKER   172    77 BLOND      FAIR        BLUE  
## 2 C-3PO           167    75 <NA>      GOLD        YELLOW  
## 3 R2-D2            96    32 <NA>      WHITE, BLUE RED  
## 4 DARTH VADER     202   136 NONE      WHITE        YELLOW  
## 5 LEIA ORGANA     150    49 BROWN     LIGHT        BROWN
```

³ This workflow (i.e. combining `mutate` and `across`) supersedes the old "scoped" variants of `mutate()` that you might have used previously

across

Alternatively, we can provide an "anonymous" **lambda function** using syntax from the **purrr** package:

- `\(x)` to indicate we're building an anonymous lambda function
- `x` the variables being passed in from `across()`

```
sw_4 <- select(starwars, name:eye_color) %>%  
  mutate(  
    across(  
      # variables to select  
      c(mass),  
      # anonymous lambda function  
      \(x) x / 100  
    )  
  )
```

across

Alternatively, we can provide an "anonymous" **lambda function** using syntax from the **purrr** package:

- `\(x)` to indicate we're building an anonymous lambda function
- `x` the variables being passed in from `across()`

```
starwars %>% select(name:eye_color) %>%  
  mutate(across(c(height, mass), \(x) x / 100)) %>%  
  head(5)
```

```
## # A tibble: 5 × 6  
##   name          height  mass hair_color skin_color eye_color  
##   <chr>         <dbl> <dbl> <chr>      <chr>      <chr>  
## 1 Luke Skywalker  1.72  0.77 blond      fair        blue  
## 2 C-3PO          1.67  0.75 <NA>      gold        yellow  
## 3 R2-D2          0.96  0.32 <NA>      white, blue red  
## 4 Darth Vader    2.02  1.36 none       white       yellow  
## 5 Leia Organa    1.5   0.49 brown      light       brown
```

5) dplyr::summarise

`summarise()` lets us **manually specify summary statistics**. It's particularly useful in combination with the `group_by()` command.

```
sw_5 <- starwars %>%  
  drop_na(height) %>%  
  group_by(species) %>%  
  summarise("Mean Height" = mean(height),  
            "Std Dev Height" = sd(height),  
            Count = n())
```

```
sw_5
```

```
## # A tibble: 38 × 4  
##   species    Mean Height Std Dev Height Count  
##   <chr>          <dbl>          <dbl> <int>  
## 1 Aleena           79             NA      1  
## 2 Besalisk        198             NA      1  
## 3 Cerean          198             NA      1  
## 4 Chagrian        196             NA      1  
## 5 Clawdite        168             NA      1  
## 6 Droid          131.          49.1      5
```

5) dplyr::summarise

Note that including `na.rm = TRUE` (or `na.rm = T`) is usually a good idea, otherwise, missing values will result in `NA`

```
## Probably not what we want
```

```
starwars %>%  
  summarise(mean_height = mean(height))
```

```
## # A tibble: 1 × 1  
##   mean_height  
##       <dbl>  
## 1          NA
```

5) dplyr::summarise

Note that including `na.rm = TRUE` (or `na.rm = T`) is usually a good idea, otherwise, missing values will result in `NA`

```
## Much better
starwars %>%
  summarise(mean_height = mean(height, na.rm = TRUE))

## # A tibble: 1 × 1
##   mean_height
##   <dbl>
## 1      175.
```

5) dplyr::summarise

The same `across()`-based workflow that we saw with `mutate` a few slides back also works with `summarise`. For example:

```
starwars %>%  
  group_by(species) %>%  
  summarise(across(where(is.numeric), mean, na.rm=T)) %>%  
  head(5)
```

```
## # A tibble: 5 × 4  
##   species height mass birth_year  
##   <chr>    <dbl> <dbl>    <dbl>  
## 1 Aleena      79     15      NaN  
## 2 Besalisk   198    102      NaN  
## 3 Cerean     198     82      92  
## 4 Chagrian   196    NaN      NaN  
## 5 Clawdite   168     55      NaN
```

5) dplyr::summarise

We can also specify **multiple summary functions** and **custom suffixes** by adding a **list**:

```
starwars %>%  
  group_by(species) %>%  
  summarise(across(where(is.numeric),  
                    list(Avg = mean, SD = sd),  
                    na.rm=T)) %>% head(5)
```

```
## # A tibble: 5 × 7  
##   species height_Avg height_SD mass_Avg mass_SD birth_year_Avg  
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>  
## 1 Aleena      79         NA         15         NA         NaN  
## 2 Besalisk    198         NA        102         NA         NaN  
## 3 Cerean      198         NA         82         NA         92  
## 4 ...  
## 5 ...
```

Other dplyr Goodies:

`group_by()` and `ungroup()`: For (un)grouping.

- Particularly useful with the `summarise()` and `mutate()` commands, as we've already seen.

```
starwars %>%  
  group_by(species) %>%  
  mutate(species_mass = mean(mass, na.rm = T),  
         species_mass_diff = mass - species_mass) %>%  
  select(name, starts_with("species")) %>%  
  ungroup() %>% head()
```

```
## # A tibble: 6 × 4  
##   name          species species_mass species_mass_diff  
##   <chr>         <chr>         <dbl>         <dbl>  
## 1 Luke Skywalker Human           81.3           -4.31  
## 2 C-3PO         Droid           69.8            5.25  
## 3 R2-D2         Droid           69.8           -37.8  
## 4 Darth Vader   Human           81.3            54.7  
## 5 Leia Organa   Human           81.3           -32.3
```

Other dplyr Goodies: slice

`slice()`: Subset rows by position rather than filtering by values.

```
starwars %>%  
  slice(c(1,5))
```

```
## # A tibble: 2 × 16  
##   name      height  mass hair_color skin_color eye_color birth_year sex  
gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>  
<chr> <chr>  
## 1 Luke Sky...   172    77 blond      fair       blue        19  
male  mascu...  
## 2 Leia Org...   150    49 brown      light      brown        19  
fema... femin...  
## # i 7 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>, skywalker <lgl>, not_skywalker  
<lgl>
```

Other dplyr Goodies: slice

Several other case-specific `slice` versions exist:

Function	Role
<code>slice_head/tail(df, n, prop, by)</code>	Grab the first <code>n / prop</code> top or bottom rows
<code>slice_min/max(df, order_by, ...)</code>	Grab the <code>n / prop</code> min or max value rows
<code>slice_sample(df, replace = F, ...)</code>	Grab random sample of rows

Other dplyr Goodies: pull

`pull()`: Extracts a column from a data frame as a vector or scalar.

- grab by name or position (positive integer, L to R)

```
starwars %>%  
  filter(gender="feminine") %>%  
  pull(height)
```

```
## [1] 150 165 150 185 163 178 184 170 166 165 168 213 167 96 178 NA  
NA
```

Exactly like using `$`, but works with pipes!

Other dplyr Goodies: count and distinct

`count()` and `distinct()`: Number and isolate unique observations.

`count()`: number of observations.

`distinct()`: first unique observations.

```
starwars %>% count(species)
```

```
## # A tibble: 38 × 2
##   species      n
##   <chr>    <int>
## 1 Aleena      1
## 2 Besalisk    1
## 3 Cerean      1
## 4 Chagrian    1
## 5 Clawdite    1
## 6 Droid       6
## 7 Dug         1
## 8 Ewok        1
## 9 Geonosian   1
## 10 Gungan     3
```

```
starwars ← mutate(starwars,
                   hair_color =
                     if_else(name == "Luke Skywalker",
                             "blonde", hair_color))
```

```
starwars %>% distinct(species)
```

```
## # A tibble: 38 × 1
##   species
##   <chr>
## 1 Human
## 2 Droid
## 3 Wookiee
## 4 Rodian
```

distinct

By default, `distinct()` **drops all other variables**.

If we're in a context where it makes sense to keep them (i.e. removing a few duplicated rows in a panel dataset), set `.keep_all = T`:

```
starwars %>% distinct(species, .keep_all = T) %>% dim()

## [1] 38 16

starwars %>% distinct(species, .keep_all = T) %>% head()

## # A tibble: 6 × 16
##   name      height  mass hair_color skin_color eye_color birth_year sex
gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl>
<chr> <chr>
## 1 Luke Sky...   172    77 blonde    fair       blue        19
male  mascu...
## 2 C-3PO        167    75 <NA>      gold       yellow      112 64 / 110
```

Other dplyr Goodies: count and distinct

If we want to add the count as a variable to the full dataframe, we can use `group_by()`, `mutate()`, and `n()`:

```
starwars %>% group_by(species) %>% mutate(num = n()) %>%  
  select(name, species, num)
```

```
## # A tibble: 87 × 3  
## # Groups:   species [38]  
##   name                species    num  
##   <chr>              <chr>    <int>  
## 1 Luke Skywalker     Human      35  
## 2 C-3PO              Droid       6  
## 3 R2-D2              Droid       6  
## 4 Darth Vader        Human      35  
## 5 Leia Organa         Human      35  
## 6 Owen Lars           Human      35  
## 7 Beru Whitesun Lars Human      35  
## 8 R5-D4              Droid       6  
## 9 Biggs Darklighter  Human      35  
## 10 Obi-Wan Kenobi     Human      35
```

Other dplyr Goodies: window functions

There is also a **whole class of window functions** for getting leads and lags, ranking, creating cumulative aggregates, etc.

- Generate leads and lags: `lag(species, 5)`, `lead(height)`
- Create rankings: `row_number()`, `min_rank()`, `dense_rank()`, `cume_dist()`, `percent_rank()`, `ntile()`
- Build cumulative aggregates: `cumsum()`, `cummin()`, `cummax()`, `cumall()`, `cumany()`, `cummean()`
- See `vignette("window-functions")`.

The final set of dplyr "goodies" are the family of **join operations**. However, these are important enough that I want to go over these concepts in a bit more depth...

- We will encounter and practice these many more times as the course progresses.

Joins

One of the mainstays of the dplyr package is merging data with the family [join operations](#).

- `inner_join(df1, df2)`
- `left_join(df1, df2)`
- `right_join(df1, df2)`
- `full_join(df1, df2)`
- `semi_join(df1, df2)`
- `anti_join(df1, df2)`

(You can find some helpful visual depictions of the different join operations [here](#).)

Joins

For our join examples, we'll use some data sets that come bundled with the nycflights13 package.

- Load it now and then inspect these data frames⁴ in your own console.

⁴ These datasets are technically stored as tibbles, which are an opinionated, modern version of data frames. For our uses we can treat them essentially interchangeably, or forcibly go between types with `as.data.frame()/as.tibble()`

Joins: Example Datasets

The `flights` dataset contains information on all flights that departed NYC in 2013:

```
head(flights)
```

```
## # A tibble: 6 × 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time  
sched_arr_time
```

```
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>  
<int>
```

```
## 1  2013     1     1     517           515         2     830  
819
```

```
## 2  2013     1     1     533           529         4     850  
830
```

```
## 3  2013     1     1     542           540         2     923  
850
```

```
## 4  2013     1     1     544           545        -1    1004  
1022
```

```
## 5  2013     1     1     554           600        -6     812  
837
```

Joins: Example Datasets

The `planes` dataset contains metadata for all plane tailnumbers within the FAA aircraft registry

```
head(planes)
```

```
## # A tibble: 6 × 9
##   tailnum year type      manufacturer model engines seats
##   <chr>   <int> <chr>      <chr>          <chr>   <int> <int>
##   <int> <chr>
## 1 N10156   2004 Fixed wing multi ... EMBRAER      EMB-...     2     55
NA Turbo...
## 2 N102UW   1998 Fixed wing multi ... AIRBUS INDU... A320...     2    182
NA Turbo...
## 3 N103US   1999 Fixed wing multi ... AIRBUS INDU... A320...     2    182
NA Turbo...
## 4 N104UW   1999 Fixed wing multi ... AIRBUS INDU... A320...     2    182
NA Turbo...
## 5 N10575   2002 Fixed wing multi ... EMBRAER      EMB-...     2     55
NA Turbo...
```

Joins: Example Datasets

The `weather` dataset contains info on the prevailing NYC weather at every minute

```
head(weather)
```

```
## # A tibble: 6 × 15
```

```
##   origin  year month   day  hour  temp  dewp humid wind_dir wind_speed  
wind_gust
```

```
##   <chr>   <int> <int> <int> <int> <dbl> <dbl> <dbl>    <dbl>    <dbl>  
<dbl>
```

```
## 1 EWR      2013     1     1     1  39.0  26.1  59.4      270      10.4
```

```
NA
```

```
## 2 EWR      2013     1     1     2  39.0  27.0  61.6      250      8.06
```

```
NA
```

```
## 3 EWR      2013     1     1     3  39.0  28.0  64.4      240      11.5
```

```
NA
```

```
## 4 EWR      2013     1     1     4  39.9  28.0  62.2      250      12.7
```

```
NA
```

```
## 5 EWR      2013     1     1     5  39.0  28.0  64.4      260      12.7
```

```
NA
```

Joins

Let's perform a **left join** to bring variables from the planes dataset into the flights dataset.

- `left_join(df1, df2)` keeps all rows of `df1`, adds variables from `df2`
- **Note:** I'm subsetting columns, but only for the sake of slide legibility

```
left_join(flights, planes) %>%  
  select(year:dep_time, carrier, flight, tailnum, type, model, engine)
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time carrier flight tailnum type  model engine  
##   <int> <int> <int>   <int> <chr>   <int> <chr>   <chr> <chr> <chr>  
## 1  2013     1     1     517 UA      1545 N14228 <NA> <NA> <NA>  
## 2  2013     1     1     533 UA      1714 N24211 <NA> <NA> <NA>  
## 3  2013     1     1     542 AA      1141 N619AA <NA> <NA> <NA>  
## 4  2013     1     1     544 B6        725 N804JB <NA> <NA> <NA>  
## 5  2013     1     1     554 DL        461 N668DN <NA> <NA> <NA>  
## 6  2013     1     1     554 UA      1696 N39463 <NA> <NA> <NA>  
## 7  2013     1     1     555 B6        507 N516JB <NA> <NA> <NA>  
## 8  2013     1     1     557 EV      5708 N820AS <NA> <NA> <NA>
```

Joins

Note that dplyr made a **reasonable guess** about which columns to join on (i.e. columns that **share the same name**). It also told us its choices:

```
## Joining, by = c("year", "tailnum")
```

However, there's an obvious problem here: the variable "year" does not have a consistent meaning across our joining datasets!

- In one it refers to the **year of flight**,
- In the other it refers to **year of construction**

Luckily, there's an easy way to avoid this problem.

- See if you can figure it out before turning to the next slide.
- Try `?dplyr::join`.

Joins: by

Solution: state explicitly which variables to join on by using the `by` argument.

- You can also rename any ambiguous columns to avoid confusion

```
left_join(
  flights,
  planes %>% rename(year_built = year),
  by = "tailnum" ## Be specific about the joining column
) %>%
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum,
  year_built, type, model) %>%
  head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 11
##   year month   day dep_time arr_time carrier flight tailnum year_built
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>      <int>
##   <chr>
## 1  2013     1     1     517     830 UA      1545 N14228     1999
```

Joins: Name Conflicts

Note what happens if we again specify the join column... but don't rename the ambiguous "year" column in at least one of the given data frames.

```
left_join(
  flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum"
) %>%
  select(contains("year"), month, day, dep_time, arr_time, carrier,
  flight, tailnum, type, model) %>%
  head(3)
```

```
## # A tibble: 3 × 11
##   year.x year.y month    day dep_time arr_time carrier flight tailnum
type  model
##   <int>  <int> <int> <int>    <int>    <int> <chr>    <int> <chr>
<chr> <chr>
## 1   2013   1999     1     1      517      830 UA        1545 N14228
Fixe... 737-...
## 2   2013   1998     1     1      533      850 UA        1714 N24211
```

Joining on Multiple Columns

Often we need to join on **multiple variables** (i.e. unit and time for panels).

Two main ways to use `by` when merging on multiple columns:

1. Rename matching columns before merging to have the same names
2. Specify columns with different names to match on with

```
by = c("xvar1" = "yvar1", "xvar2" = "yvar2", ... )
```

Joining on Multiple Columns

To see these, let's get info from the `weather` dataset:

```
flights <- flights
weather_sub <- select(weather, year, month, day, hour, temp, humid,
  starts_with("wind"))
```

This dataset contains info on the temperature, humidity, and wind conditions at each hour of the day in NYC during 2013 - useful information for understanding reasons for flight delays!

Joining on Multiple Columns

Suppose we want to have an approximation of the weather conditions before each flight. Since weather is only to the nearest hour, let's round flight departure to the closest hour⁵:

```
flights <- mutate(flights,  
                  dep_hr = case_when(  
                    nchar(dep_time) == 3 ~ substr(dep_time,1,1) %>%  
as.numeric(),  
                    nchar(dep_time) == 4 ~ substr(dep_time,1,2) %>%  
as.numeric(),  
                    TRUE ~ as.numeric(NA)  
                  )  
)
```

⁵ I'm doing this to get to a starting point of different variable names - in reality we could just jump straight to merging on the same names here.

Joins: by (Renaming First)

Here we want to join on time (year, month, and departure hour).

We could begin by renaming `hour` in the weather dataset to match the flights data:

```
left_join(  
  flights,  
  weather_sub %>% rename(dep_hr = hour), ## Rename to match  
  by = c("year", "month", "day", "dep_hr") ## Specify join columns  
) %>%  
  select(year, month, day, dep_hr, flight, temp, humid) %>%  
  head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 7  
##   year month   day dep_hr flight  temp humid  
##   <int> <int> <int>   <dbl>   <int> <dbl> <dbl>  
## 1  2013     1     1       5    1545  39.0  64.4  
## 2  2013     1     1       5    1545  39.0  61.6  
## 3  2013     1     1       5    1545  39.9  54.8
```

Joins: by (Merging on Different Names)

Alternatively, we could perform the same join without renaming (R will keep the X data's variable name for any naming differences)

```
left_join(  
  flights,  
  weather_sub,  
  by = c("year" = "year", "month" = "month", "day" = "day", "dep_hr" =  
"hour") ## Be specific about the joining column ("Xname" = "Yname")  
) %>%  
  select(year, month, day, dep_hr, flight, temp, humid) %>%  
  head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 7  
##   year month   day dep_hr flight  temp humid  
##   <int> <int> <int>   <dbl>   <int> <dbl> <dbl>  
## 1  2013     1     1       5    1545  39.0  64.4  
## 2  2013     1     1       5    1545  39.0  61.6  
## 3  2013     1     1       5    1545  39.9  54.8
```

Mutating Joins

left joins are probably the most common join we'll do, but we can perform a wide range of **mutating joins** with other join functions:

Join Function	Description
<code>left_join(df1, df2)</code>	Add variables from <code>df2</code> into <code>df1</code> (keep all rows of <code>df1</code>)
<code>right_join(df1, df2)</code>	Add variables from <code>df1</code> into <code>df2</code> (keep all rows of <code>df2</code>)
<code>full_join(df1, df2)</code>	Combine <code>df1</code> and <code>df2</code> (keep all rows of <code>df1</code> and <code>df2</code>)
<code>inner_join(df1, df2)</code>	Keep only observations from <code>df1</code> with matches in <code>df2</code>

Cross Join

A somewhat different mutating join is the **cross join**.

`cross_join(df1, df2)` matches each row in `df1` with each row in `df2`

- Results in a `nrow(df1) * nrow(df2)` row dataframe
- Ignores the data in the columns
 - i.e. doesn't have a `by` argument
- Useful for expanding an object with all possible combinations of values

join_by()

Another helpful tool for conducting **conditional joins** is `join_by()`

- Equality (`=`) or inequality (`<`, `>`, `≤`, `≥`, `≠`)
- Rolling (`closest()`)
- Overlap (`between()`, `within()`, `overlaps()`)
- Place in the `by` argument

join_by()

Example: add to weather the info for the first flight that occurred each hour into the weather data

```
# get all flights on 9/10, convert scheduled departure time to combined numeric
flights_9_10 <- filter(flights, month = 9 & day = 10) %>%
  mutate(dep_num = hour + minute/60) %>%
  select(-hour)

weather_close <- left_join(weather_sub, flights_9_10,
  by = join_by(year, month, day, # variables to join on
               closest(hour ≤ dep_num)          )) %>% #
condition to join on
  arrange(tailnum) # bring the matches to the top
```

join_by()

Note that we get **duplicated matches**.

Q: Why?

join_by()

A: weather isn't unique on the matching variables!

```
# group on the hour of each day
group_by(weather_sub, year, month, day, hour) %>%
  # get only the 2nd or higher entry by hour of each day
  filter(row_number() > 1) %>%
  # how many rows are duplicates?
  nrow()
```

```
## [1] 17402
```

join_by()

Get the first weather observation each hour, then rejoin

```
weather_dist_9_10 <- distinct(weather_sub, year, month, day, hour,  
  .keep_all = T) %>%  
  filter(month = 9 & day = 10)  
  
weather_close <-  
  left_join(weather_dist_9_10, flights_9_10,  
    by = join_by(year, month, day, # variables to join on  
      closest(hour ≤ dep_num)    )) # condition  
to join on
```

join_by()

Dang, we *still* have too many rows

- `weather_dist_9_10` has 24 rows
- `weather_close` has 179!

Q: Why?

join_by()

A: multiple flights share the same scheduled departure time!

```
group_by(flights_9_10, dep_num) %>%  
  filter(row_number() > 1) %>%  
  # how many rows are duplicates?  
  nrow()
```

```
## [1] 619
```

Important lesson: always

1. Understand/check for uniqueness *before* merging
2. Make sure the rows of the returned dataframe match your expectations
3. If not, **investigate!**

Filtering Joins

We can also perform **filtering joins** to restrict samples based on matches/non-matches across datasets:

Join Function	Description
<code>semi_join(df1, df2)</code>	return all rows of <code>df1</code> with a match in <code>df2</code>
<code>anti_join(df1, df2)</code>	return all rows of <code>df1</code> without a match in <code>df2</code>

tidyr

Key tidyr verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt").⁶
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast").⁷
3. `separate_wider_delim/separate_longer_delim`: Separate (i.e. split) one column into multiple columns/multiple rows.
4. `unite`: Unite (i.e. combine) multiple columns into one.

Let's practice these verbs together in class.

- Side question: Which of `pivot_longer` vs `pivot_wider` produces "tidy" data?

⁶ Updated version of `tidyr::gather`.

⁷ Updated version of `tidyr::spread`.

1) tidyr::pivot_longer

Use `pivot_longer()` to go **from wide to long**⁸:

```
stocks <- data.frame( ## Could use "tibble" instead of "data.frame" if  
you prefer  
  time = as.Date('2009-01-01') + 0:1,  
  X = rnorm(2, 0, 1),  
  Y = rnorm(2, 0, 2),  
  Z = rnorm(2, 0, 4)  
)  
stocks
```

```
##           time           X           Y           Z  
## 1 2009-01-01 -0.14815441 -0.4698242  1.619799  
## 2 2009-01-02  0.06580843 -1.6469002 -1.578764
```

1) tidyr::pivot_longer

Use `pivot_longer()` to go **from wide to long**⁸:

```
stocks %>% pivot_longer(-time, names_to="stock", values_to="price")
```

```
## # A tibble: 6 × 3
##   time      stock price
##   <date>    <chr>  <dbl>
## 1 2009-01-01 X      -0.148
## 2 2009-01-01 Y      -0.470
## 3 2009-01-01 Z        1.62
## 4 2009-01-02 X        0.0658
## 5 2009-01-02 Y      -1.65
## 6 2009-01-02 Z      -1.58
```

⁸ Note that both pivot functions have a lot of handy options for modifying names.

1) tidyr::pivot_longer

We could also manually specify the columns to pivot (useful when we want to pivot on just a subset of columns)

```
stocks %>% pivot_longer(cols = c(X, Y, Z), names_to="stock",  
  values_to="price")
```

```
## # A tibble: 6 × 3  
##   time      stock price  
##   <date>    <chr>  <dbl>  
## 1 2009-01-01 X      -0.148  
## 2 2009-01-01 Y      -0.470  
## 3 2009-01-01 Z        1.62  
## 4 2009-01-02 X       0.0658  
## 5 2009-01-02 Y      -1.65  
## 6 2009-01-02 Z      -1.58
```

1) tidyr::pivot_longer

Let's quickly save the "tidy" (i.e. long) stocks data frame for use on the next slide.

```
## Write out the argument names this time: i.e. "names_to=" and  
"values_to="  
tidy_stocks ← pivot_longer(stocks, -time, names_to="stock",  
values_to="price")
```

2) tidyr::pivot_wider

Use `pivot_wider()` to go **from long to wide**:

```
tidy_stocks %>% pivot_wider(names_from=stock, values_from=price)
```

```
## # A tibble: 2 × 4
##   time          X          Y          Z
##   <date>      <dbl>   <dbl>   <dbl>
## 1 2009-01-01 -0.148  -0.470   1.62
## 2 2009-01-02  0.0658 -1.65   -1.58
```

```
tidy_stocks %>% pivot_wider(names_from=time, values_from=price)
```

```
## # A tibble: 3 × 3
##   stock 2009-01-01 2009-01-02
##   <chr>   <dbl>      <dbl>
## 1 X      -0.148      0.0658
## 2 Y      -0.470     -1.65
## 3 Z       1.62     -1.58
```

Note the second ex. has effectively transposed the data.

Aside: Remembering the pivot_* syntax

There's a long-running joke about no-one being able to remember Stata's "reshape" command. ([Exhibit A](#).)

It's easy to see this happening with the `pivot_*` functions too. However, I find that I never forget the commands as long as I remember the argument order is *"names"* then *"values"*.

3) Separate

tidyr has several `separate_direction_method` functions that make it easy to separate values in a column into multiple columns/rows, where

- `direction` informs whether the data spread
 - `wide` (`_wider_`), or
 - expand each cell into multiple rows (`_longer_`)
- `method` instructs the way to split a cell:
 - `delim` to split on a delimiter (i.e. "." or "/")
 - `position` to split at fixed widths
 - `regex` to split with a regular expression (i.e. `a(? ≤ d)`)

3) Separate

Let's try splitting some economists' names.

```
economists = data.frame(name = c("Adam.Smith", "Paul.Samuelson",  
"Milton.Friedman"))  
economists
```

```
##           name  
## 1   Adam.Smith  
## 2 Paul.Samuelson  
## 3 Milton.Friedman
```

3) tidyr::separate_wider_delim

To split names into two columns by splitting at the period, we can use `separate_wider_delim`:

```
economists %>% separate_wider_delim(name, # column(s) to separate
                                   delim = ".", # delimiter to split on
                                   names = c("first_name", "last_name")) # name of new
variables
```

```
## # A tibble: 3 × 2
##   first_name last_name
##   <chr>      <chr>
## 1 Adam      Smith
## 2 Paul      Samuelson
## 3 Milton    Friedman
```

3) tidyr::separate_wider_regex

If you know regular expressions, you can use `separate_wider_regex` to accomplish the same task:

```
economists %>%  
  separate_wider_regex(name,  
    patterns = c(first_name = "[:alpha:]+", ".", last_name = "  
[:alpha:]+"))
```

```
## # A tibble: 3 × 2  
##   first_name last_name  
##   <chr>      <chr>  
## 1 Adam      Smith  
## 2 Paul      Samuelson  
## 3 Milton    Friedman
```

3) tidyr::separate_longer_delim

A related function is `separate_longer_delim`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data).

Let's see its use with some occupation data

```
jobs <- data.frame(  
  name = c("Jack", "Jill"),  
  occupation = c("Homemaker", "Philosopher, Philanthropist,  
  Troublemaker")  
)  
jobs
```

```
##      name                occupation  
## 1 Jack                Homemaker  
## 2 Jill Philosopher, Philanthropist, Troublemaker
```

3) tidyr::separate_longer_delim

We can expand the data to have one row for each name and occupation combination:

```
## Now split out Jill's various occupations into different rows  
jobs %>% separate_longer_delim(occupation, delim = ", ")
```

```
##   name      occupation  
## 1 Jack      Homemaker  
## 2 Jill      Philosopher  
## 3 Jill Philanthropist  
## 4 Jill      Troublemaker
```

4) tidyr::unite

`unite()` allows us to **collapse multiple columns into a single column**

Suppose we have daily small business revenues:

```
rev <- data.frame(
  year = rep(2016, times = 4),
  month = rep(1, times = 4),
  day = 1:4,
  revenue = rnorm(4, mean = 100, sd = 10)
)
rev
```

```
##   year month day  revenue
## 1 2016     1   1 101.86306
## 2 2016     1   2  84.65713
## 3 2016     1   3  99.75226
## 4 2016     1   4  98.66922
```

4) tidyr::unite

We can use `unite` to combine the three date components into a single character column⁹:

```
## Combine "yr", "mnth", and "dy" into one "date" column
rev_u <- rev %>% unite(col = date, # name of new column
                      c("month", "day", "year", ), # columns to unite
                      sep = "-") # separator to use
```

```
rev_u
```

```
##      date    revenue
## 1 1-1-2016 101.86306
## 2 1-2-2016  84.65713
## 3 1-3-2016  99.75226
## 4 1-4-2016  98.66922
```

⁹ Set the argument `remove = T` to keep the original input columns

4) tidyr::unite

If we want to convert the new character column to another type (e.g. date or numeric) then you will need to modify it using `mutate`.

For example, we can use the [lubridate](#) package's super helpful date conversion functions to convert our new variable to a date:

```
pacman::p_load(lubridate)
rev_u <- mutate(rev_u,
                date = mdy(date))
class(rev_u$date)
```

```
## [1] "Date"
```

Other tidyr goodies

Use `crossing` to get the full combination of a group of variables.¹⁰

```
crossing(side=c("left", "center", "right"),  
         height=c("top", "middle", "bottom"))
```

```
## # A tibble: 9 × 2  
##   side    height  
##   <chr>  <chr>  
## 1 center bottom  
## 2 center middle  
## 3 center top  
## 4 left   bottom  
## 5 left   middle  
## 6 left   top  
## 7 right  bottom  
## 8 right  middle  
## 9 right  top
```

¹⁰ Base R alternative: `expand.grid`.

Other tidyr goodies

Use `crossing` to get the full combination of a group of variables.¹¹

```
crossing(side=c("left", "center", "right"),  
         height=c("top", "middle", "bottom"))
```

See `?expand` and `?complete` for more specialized functions that allow you to fill in (implicit) missing data or variable combinations in existing data frames.

¹¹ Base R alternative: `expand.grid`.

Table of Contents

1. Prologue
2. Tidyverse Overview
3. Pipes
4. Data Wrangling with dplyr
 - filter, arrange, select, mutate, summarise, joins
5. Data Tidying with tidyr
 - pivot longer / pivot wider
 - separate
 - unite